

Sentinel-Ops — Full Operational & Architectural Report

Hybrid Edition: Historical Overview + Detailed Log of Today's Work

Generated: 14 February 2026

1. Executive Summary

This document provides a complete, end-to-end overview of the Sentinel-Ops project as of 14 February 2026, including:

- Repository architecture
- RAG ingestion pipeline
- IaC structure
- YAML + Python components
- Windows ↔ Ubuntu DevOps VM ↔ GitHub linkage
- Today's operational progress
- Debugging steps
- Status index
- Next actions

This report is designed as a **canonical reference** for both human operators and the Sentinel-Ops RAG engine.

2. Project Context & Objectives

Sentinel-Ops is a multi-LLM, anti-drift, version-controlled operational knowledge system designed to:

- Maintain a curated **canon** of authoritative documents
- Provide reproducible **infrastructure-as-code**
- Support **RAG-based reasoning** across the entire project
- Ensure continuity across Windows, Ubuntu DevOps VM (VM 100), and GitHub
- Enable long-term operational resilience

The system is built around:

- **GitHub** as the source of truth
- **Windows 11** as the primary workstation
- **Ubuntu DevOps VM 100** as the execution environment

- **Ollama** for local LLM inference
 - **ChromaDB** for vector storage
 - **Python** for ingestion + query pipelines
-

3. Repository Architecture (Current)

```
sentinel-ops/
    └── canon/                      # Curated, authoritative documents
        ├── Architecture.md
        ├── Infrastructure-as-Code (IaC) & Security Baseline Plan.md
        ├── Roadmaps/
        └── Policies/

    └── infra/                      # Infrastructure-as-Code
        ├── terraform/
        ├── kubernetes/
        ├── ansible/
        └── artifacts/
            └── sboms/
                └── host-ubuntu-24.04.sbom.json (50MB)

    └── reports/                     # Generated reports (PDF-ready MD)

    └── rag/                         # RAG engine
        ├── ingest.py
        ├── query.py
        ├── config.yaml
        └── vectorstore/

└── README.md
```

4. Cross-System Mapping

4.1 Windows File Structure

Primary working directory:

C:\Users\Admin\Documents\sentinel-ops\

RAG engine:

C:\Users\Admin\Documents\sentinel-ops\rag\

Ollama models stored under:

C:\Users\Admin\.ollama\models\

Python environment:

C:\Users\Admin\AppData\Local\Programs\Python\Python312\

4.2 Ubuntu DevOps VM (VM 100)

Mounted repo location:

/home/andy/sentinel-ops/

Used for:

- Terraform execution
 - Kubernetes manifests
 - CI/CD testing
 - SBOM generation
 - DevOps automation
-

4.3 GitHub Integration

Remote origin:

<https://github.com/<your-org>/sentinel-ops.git>

Windows ↔ GitHub:

- Git CLI
- VS Code
- GitHub Desktop

Ubuntu VM ↔ GitHub:

- SSH keys stored under /home/andy/.ssh/
 - Git operations executed inside VM 100
-

5. IaC Overview

5.1 Terraform

Located under:

infra/terraform/

Typical structure:

```
main.tf  
variables.tf  
outputs.tf  
providers.tf  
modules/
```

Used for:

- VM provisioning
 - Network configuration
 - Storage
 - Security baselines
-

5.2 Kubernetes Manifests

Located under:

```
infra/kubernetes/
```

Includes:

- Deployments
 - Services
 - Ingress
 - ConfigMaps
 - Secrets (referenced, not stored)
-

5.3 Ansible

Located under:

```
infra/ansible/
```

Used for:

- Host configuration
 - Package installation
 - Hardening
 - SBOM generation tasks
-

6. RAG Engine Components

6.1 config.yaml

Example structure:

```
include_paths:
  - "../canon"
  - "../infra"
  - "../reports"

model: "gemma3:4b"

vectorstore_path: "./vectorstore"

chunk_size: 500
chunk_overlap: 50
```

6.2 ingest.py (Corrected Version)

Key responsibilities:

- Walk repo
- Read files
- Chunk text
- Embed chunks using `nomic-embed-text`
- Store vectors in ChromaDB

Important notes:

- SBOMs are extremely large
 - Chunking prevents context overflow
 - Embedding is CPU-bound
 - Ingestion is linear and deterministic
-

6.3 query.py

Responsibilities:

- Accept user question
- Retrieve relevant chunks
- Build prompt
- Send to Gemma
- Return answer

Gemma is used **only** for generation, not embeddings.

7. Today's Operational Log (14 Feb 2026)

7.1 Tasks Completed

Index	Task	Status	Notes
1	Validated repo structure	✓	canon/, infra/, rag/ confirmed
2	Installed ChromaDB	✓	Python environment stable
3	Installed Ollama models	✓	gemma3:4b + nomic-embed-text
4	Debugged ingestion errors	✓	Context-length errors resolved
5	Implemented safe chunking	✓	Prevents infinite loops
6	Rewrote ingest.py	✓	Stable, deterministic
7	Identified SBOM bottleneck	✓	50MB file causing long runtime
8	Confirmed ingestion progress	✓	Verified via Task Manager
9	Produced canonical report	✓	This document

7.2 Tasks In Progress

Task	Status	Notes
Ingestion of SBOM	⌚	Very large file; CPU-bound
Full vectorstore build	⌚	Continues after SBOM
RAG validation	⌚	Pending ingestion completion

7.3 Tasks Blocked

Task	Blocked By	Notes
Query testing	Ingestion	Must complete vectorstore first
RAG tuning	Ingestion	Requires full corpus

8. Performance Observations

Task Manager shows:

- Ollama ~40% CPU
- Python ~10% CPU
- Disk I/O active
- Memory stable

This is the expected signature of:

- active chunking
- active embedding
- active vectorstore writes

No signs of:

- infinite loops
 - memory leaks
 - deadlocks
-

9. Recommendations

9.1 Exclude SBOMs from RAG

SBOMs:

- add noise
- slow ingestion
- degrade retrieval
- provide no semantic value

Recommendation:

```
if file.endswith(".sbom.json"):  
    continue
```

9.2 Add File Size Guard

```
if os.path.getsize(full_path) > 5_000_000:  
    continue
```

9.3 Add Logging

- per-file
 - per-chunk
 - timing
-

10. Next Steps

1. Allow ingestion to finish or skip SBOMs

2. Run `query.py`
 3. Validate retrieval quality
 4. Tune chunk size if needed
 5. Add ingestion logs
 6. Add RAG evaluation prompts
 7. Commit this report to `reports/`
-

11. Appendix A — Full Repo Tree Snapshot

(Generated from your description and validated structure)

```
sentinel-ops/
  └── canon/
    ├── Architecture.md
    ├── Infrastructure-as-Code (IaC) & Security Baseline Plan.md
    └── ...
  └── infra/
    ├── terraform/
    ├── kubernetes/
    ├── ansible/
    ├── artifacts/
      └── sboms/
        └── host-ubuntu-24.04.sbom.json
  └── rag/
    ├── ingest.py
    ├── query.py
    ├── config.yaml
    └── vectorstore/
  reports/
```

12. Appendix B — RAG Pipeline Diagram (Text)

```
[canon/] ----\
[infra/] -----+--> ingest.py --> embeddings --> ChromaDB --> query.py --> Gemma -->
Answer
[reports/] --/
```

13. Closing Summary

This document now serves as:

- a canonical reference
- a RAG-ingestible artifact
- a PDF-ready operational report

- a cross-system architectural map
- a complete log of today's work

It captures the full historical context **and** the detailed operational progress of 14 February 2026.

If you want, I can also generate a **short executive summary version** or a **diagram-focused version** to accompany this main report.