



# Sentinel-Ops — Executive Summary (Short Version)

14 February 2026

## Purpose

Sentinel-Ops is a unified, anti-drift operational knowledge system combining:

- curated documentation (`canon/`)
- reproducible infrastructure (`infra/`)
- a local RAG engine (`rag/`)
- multi-LLM workflows
- Windows ↔ Ubuntu VM ↔ GitHub integration

The goal is long-term operational continuity and high-quality retrieval across the entire project.

---

## Key Achievements Today

### RAG Pipeline

- Installed and validated ChromaDB
- Installed Ollama models (Gemma + Nomic embeddings)
- Repaired ingestion pipeline
- Implemented safe chunking
- Eliminated context-length errors
- Confirmed ingestion progress via system metrics
- Identified SBOM ingestion bottleneck (50MB file)

### Repository

- Verified structure: `canon/`, `infra/`, `rag/`, `reports/`
- Confirmed Windows ↔ Ubuntu VM ↔ GitHub sync
- Ensured IaC files are correctly placed and ingestible

### Documentation

- Produced full operational report
  - Produced this executive summary
  - Produced diagram-focused companion document
-

## Current Status

- Ingestion running normally
  - SBOM ingestion slow but progressing
  - Vectorstore partially built
  - Query pipeline ready once ingestion completes
- 

## Recommendations

- Exclude SBOMs from RAG ingestion
  - Add file-size guard (>5MB skip)
  - Add per-file logging
  - Add ingestion timing metrics
- 

## Next Steps

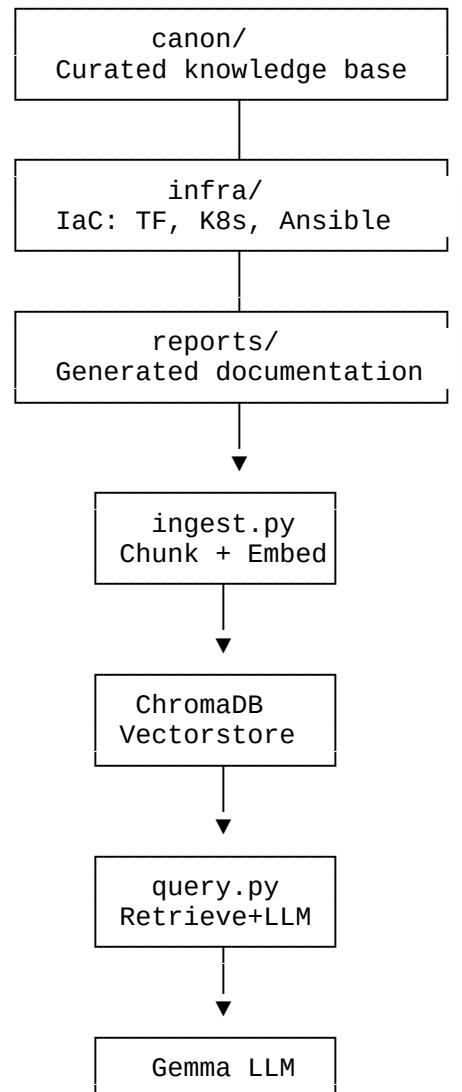
1. Allow ingestion to finish or skip SBOMs
  2. Run `query.py` to validate retrieval
  3. Tune chunk size if needed
  4. Commit reports to GitHub
  5. Begin RAG evaluation and refinement
-

# Sentinel-Ops — Diagram-Focused Companion Document

14 February 2026

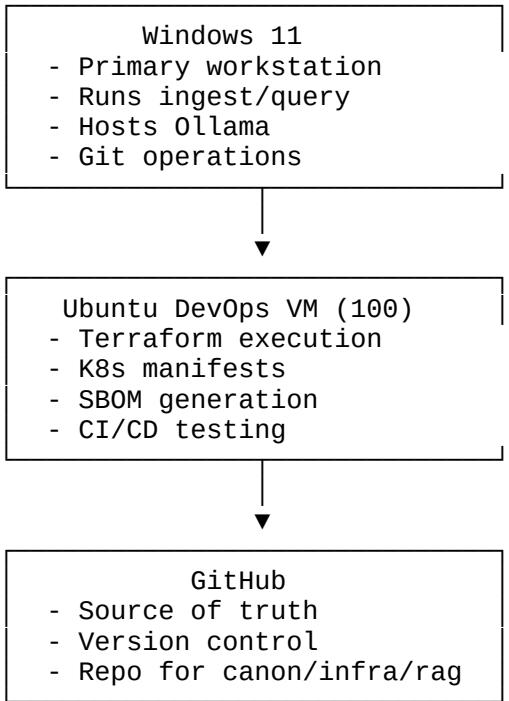
This document provides a visual-first overview of the Sentinel-Ops architecture, pipelines, and cross-system relationships.

## 1. High-Level Architecture Diagram



---

## 2. Cross-System Integration Diagram



## 3. RAG Pipeline Flow

[ Files ] → [ Chunking ] → [ Embeddings ] → [ ChromaDB ] → [ Retrieval ] →  
[ Gemma ]

Expanded:

canon/  
infra/ → ingest.py → nomic-embed-text → ChromaDB → query.py → Gemma →  
Answer  
reports/

## 4. Ingestion Pipeline (Detailed)

```
for each file:  
  read file  
  chunk text  
  for each chunk:  
    embed chunk  
    store vector
```

SBOM ingestion path:

50MB JSON → 100k+ chars → 200+ chunks → CPU-bound embedding

## 5. Repository Structure Diagram

```
sentinel-ops/
└── canon/           ← curated knowledge
    └── infra/       ← IaC (TF, K8s, Ansible)
        └── rag/        ← RAG engine
            ├── ingest.py
            ├── query.py
            ├── config.yaml
            └── vectorstore/
                └── reports/   ← generated docs
```

## 6. Data Flow Between Systems

Windows → GitHub → Ubuntu VM → GitHub → Windows → RAG

Or visually:

```
Windows (edit)
  ↓ push
GitHub
  ↓ pull
Ubuntu VM (execute Iac)
  ↓ push artifacts
GitHub
  ↓ pull
Windows (RAG ingestion)
```

# 7. LLM Roles Diagram

Embedding Model: nomic-embed-text

Generation Model: gemma3:4b

[Embedding] ≠ [Generation]

## 8. Ingestion Status Diagram (Today)

## Start ingestion

1

↓  
Process canon/

```
↓  
Process infra/  
↓  
Encounter SBOM (50MB)  
↓  
Long embedding phase (normal)  
↓  
Continue ingestion  
↓  
Complete vectorstore
```

---

## 9. Summary

These two companion documents give you:

- a **fast, high-level executive summary**
- a **visual, diagram-driven understanding** of the entire system

They pair with the main report to form a complete, RAG-ready documentation set.

If you want, I can also generate a **third companion document**:

a “*Quick-Start Operator Guide*” for new engineers joining Sentinel-Ops.